

Dear Hiring Team,

Building AI infrastructure for India's languages and population scale is one of the most important engineering challenges of this decade. Sarvam is the team best positioned to solve it, and I want to be part of that effort.

My experience brings two things that are directly relevant to Sarvam's work:

Production RAG and LLM infrastructure. At FrontPage, I built multiple RAG-based chatbots using OpenAI APIs -- a support bot, a learning assistant, and a general assistant. The moderation bot alone reduced human intervention in support queries by 90%. Beyond chatbots, I integrated vector-powered personalized feeds using BERT embeddings, GPT calls, and Milvus as a vector database, increasing engagement by 40%. I also built the underlying data pipeline using real-time events from BigQuery to construct richer user vectors. This is not surface-level prompt engineering -- I have built the full stack from embedding pipelines to vector retrieval to serving layer, I understand the engineering tradeoffs involved in making these systems reliable and fast in production.

High-throughput backend and data pipeline engineering. AI infrastructure is ultimately backend infrastructure with unique constraints. I built a real-time data pipeline in Go with Kafka and ProtoBuffers that scaled from 15k to 200k DAU in 14 weeks. I have experience with Clickhouse for analytical workloads, Redis for low-latency serving, and Temporal for reliable workflow orchestration. At VaultWealth, I migrated complex workflows to Temporal, eliminating 90% of unexpected failures. These are the same building blocks needed to build reliable LLM serving, data processing, and model evaluation pipelines.

What stands out about Sarvam is the India-first approach. Building LLM infrastructure which works for Indian languages, handles the diversity of Indian users, and operates at India's scale requires a fundamentally different engineering approach than adapting Western AI products. I want to work on problems where the engineering decisions are driven by the unique needs of the Indian market.

Happy to chat more about any of this -- especially the RAG and vector database work.

Cheers,
Sarthak